

Regression

Machine Learning for Behavioral Data
February 24, 2026

Today's Topic

Week	Lecture/Lab
1	Data Exploration
2	Regression
3	Classification + Model Evaluation
4	Time Series Prediction
5	Time Series Prediction
6	Time Series Prediction
7	Unsupervised Learning
8	Spring Break

Getting ready for today's lecture...

- Join us on SpeakUp

<https://go.epfl.ch/mlbd-speakup>

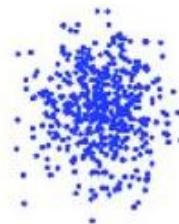


Short quiz about the past...

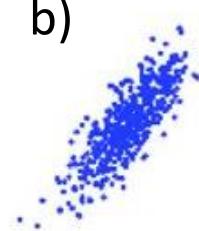
Which of the four graphs have the following properties:

High Pearson's Correlation, High Mutual Information

a)



b)

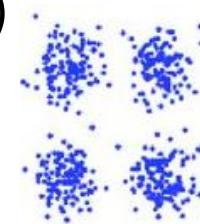


c) None

d)



e)



SpeakUp Poll!

0 votes

14/02/2022 21:26, by me

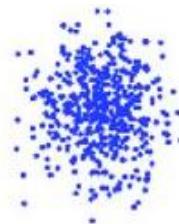
0 comments

Short quiz about the past...

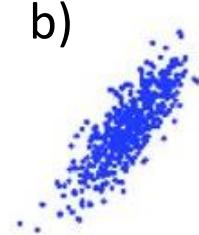
Which of the four graphs have the following properties:

High Pearson's Correlation, Low Mutual Information

a)



b)

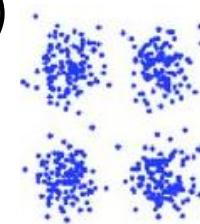


c) None

d)



e)



SpeakUp Poll!

0 votes

14/02/2022 21:26, by me

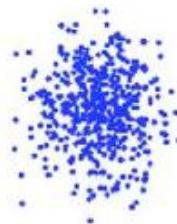
0 comments

Short quiz about the past...

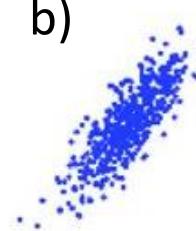
Which of the four graphs have the following properties:

Low Pearson's Correlation, Low Mutual Information

a)



b)

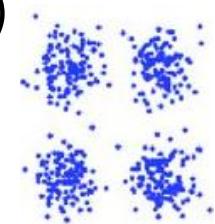


c) None

d)



e)



SpeakUp Poll!

0 votes

14/02/2022 21:26, by me

0 comments

Today's Use Case: Flipped Classroom Course

- Participants: 288 EPFL students of a course taught in *flipped classroom* mode with a duration of 10 weeks
- Structure:
 - Preparation: watch videos (and solve simple quizzes) on **new content** at home as a preparation for the lecture
 - Lecture: discuss open questions and solve more complex tasks
 - Lab session: solve paper-an-pen assignments
- Data: clickstream data (all interactions of the student with the system)

Agenda

- **Linear Regression**
- Generalized Linear Models
- Mixed-Effect Models
- Regression for Time-Series Data
- Performance Metrics

Idea

- In regression, a single aspect of the data (output variable) is modeled by some combination of other aspects of the data (input variables)

More formal

- In regression, a single aspect of the data (output variable) is modeled by some combination of other aspects of the data (input variables)
- Given: N data points (y_n, \mathbf{x}_n) , where y_n is the n^{th} output variable and \mathbf{x}_n is a D-dimensional vector of input variables
- Goal: $y_n \approx f(\mathbf{x}_n)$

Usage

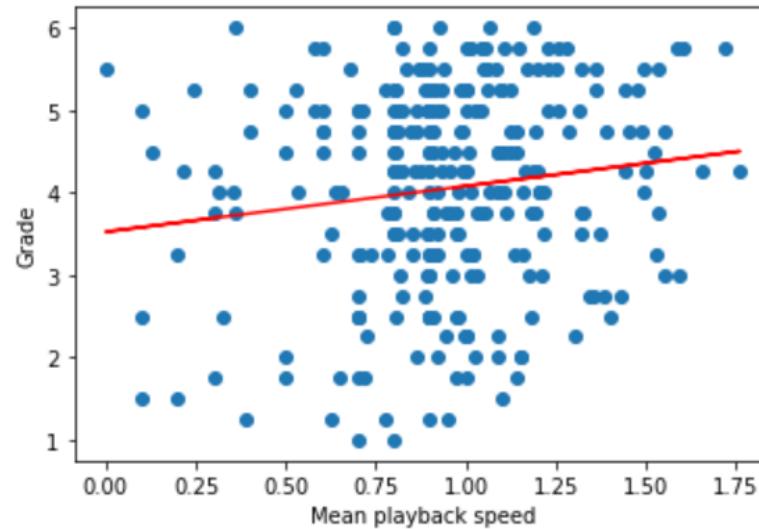
- *Prediction*: predict the output for a new (unseen) input vector x
- *Interpretation*: analyze the relationships between the variables (what effect the input variables have on the output variable)

Example | Mean playback speed

x-axis: Mean playback speed of videos

y-axis: Course grade

Each point is one student



Students who watch the videos faster tend to have better grades.

Linear Regression

The output variable y_n with $n = 1, \dots, N$ is modeled by a **linear** combination of the input variables $x_{n,d}$ with $d = 1, \dots, D$.

$$y_n = \beta_0 + \beta_1 x_{n,1} + \cdots + \beta_D x_{n,D} + \epsilon_n$$

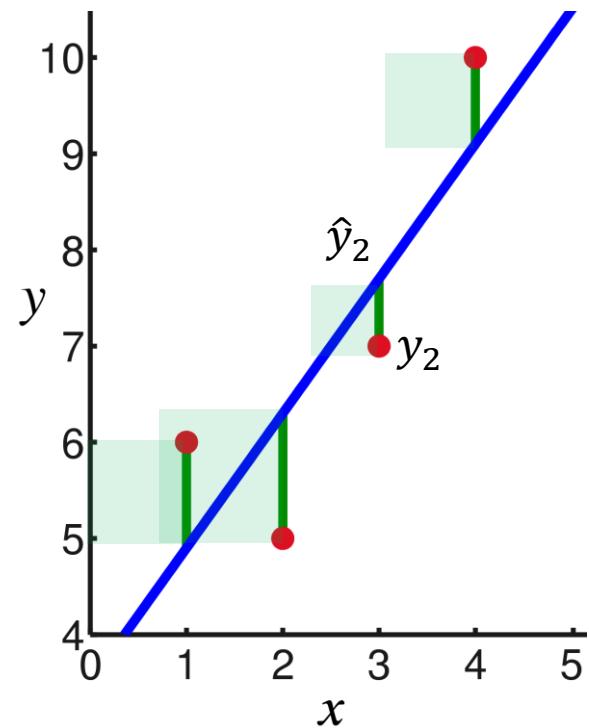
where ϵ_n are error terms that should be as small as possible and $\epsilon_n \sim N(0, \sigma^2)$.

Goal: find optimal parameters

Find parameters $\hat{\beta}$ that minimize

$$\sum_{n=1}^N (y_n - \tilde{x}_n^T \cdot \hat{\beta})^2$$

with $\tilde{x}_n = \begin{bmatrix} 1 \\ x_{n,1} \\ \dots \\ x_{n,D} \end{bmatrix}$ and $\hat{\beta} = \begin{bmatrix} \hat{\beta}_0 \\ \hat{\beta}_1 \\ \dots \\ \hat{\beta}_D \end{bmatrix}$



Fitting the parameters

$$grade = \beta_0 + \beta_1 \cdot time_in_problem + \beta_2 \cdot percentage_correct$$

Formula: grade~ch_time_in_prob_sum+wa_num_subs_perc_correct

Family: gaussian

Estimator: OLS

Std-errors: non-robust CIs: standard 95%

Inference: parametric

Number of observations: 288

R^2: 0.110

R^2_adj: 0.104

Log-likelihood: -449.516

AIC: 905.031

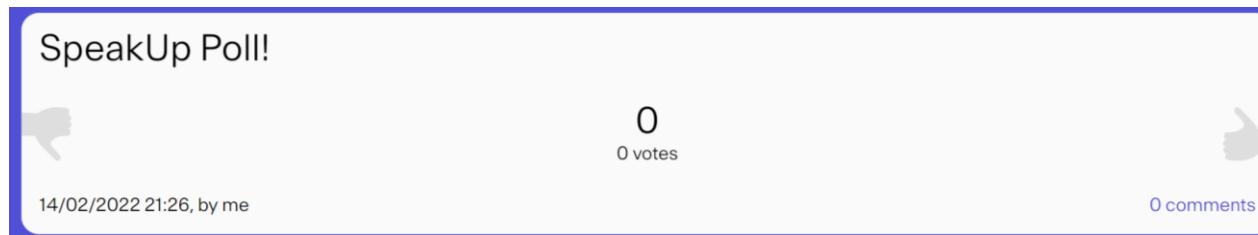
BIC: 916.020

Fixed effects:

	Estimate	2.5_ci	97.5_ci	SE	DF	T-stat	P-val	Sig
Intercept	3.410119	3.148091	3.672148	0.133123	285	25.616335	0.000000	***
ch_time_in_prob_sum	0.000157	0.000094	0.000220	0.000032	285	4.921856	0.000001	***
wa_num_subs_perc_correct	0.716132	0.035683	1.396581	0.345700	285	2.071542	0.039208	*

Influence of input variables

$$grade = 3.4 + 0.000016 \cdot time_in_problem + 0.72 \cdot percentage_correct$$



Which of the input variables has the largest impact on *grade*?

- a) *time_in_problem*
- b) *percentage_correct*
- c) I don't know

Different units of measurements

$$grade = 3.2 + 0.000016 \cdot time_in_problem + 0.72 \cdot percentage_correct$$

increase in
time_in_problem
by 1s -> increase of
grade by 0.000016

increase in
percentage_correct by 1
percentage point -> increase of
grade by 0.72

Transformation: Z-Scores

$$\tilde{x}_{n,d} = \frac{x_{n,d} - \bar{x}_d}{\sigma(x_d)}$$

$d = 1, \dots, D$

$n = 1, \dots, N$

- Standardization via z-score: $\tilde{x}_{n,d}$ denotes the distance between the raw feature $x_{n,d}$ and the sample mean \bar{x}_d (in units of the standard deviation)

Transformation: Example

$$grade = 4.05 + 0.35 \cdot time_in_problem + 0.15 \cdot percentage_correct$$

Example in Jupyter Notebook

Transformation: Summary

- Lets us compare the impact of input variables with different scales/units of measurements (e.g., time in problem in seconds and percentage correct)
- Reduces interpretability of individual input variables

Interpretation: Caveat

$$bodyfat = -45.95 + 0,99 \cdot abdomen - 0,33 \cdot weight$$

Can I conclude that heavier people (higher weight) have a lower bodyfat percentage?

Interpretation: Caveat

$$bodyfat = -45.95 + 0,99 \cdot abdomen - 0,33 \cdot weight$$

Can I conclude that heavier people (higher weight) have a lower bodyfat percentage?

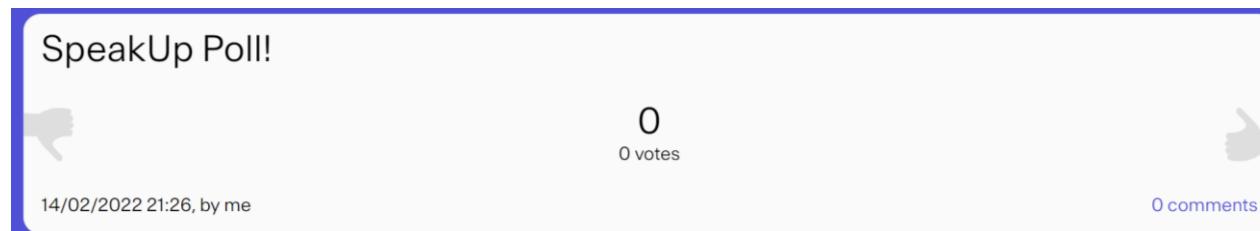
- a) Yes
- b) No
- c) I don't know

SpeakUp Poll!

0
0 votes

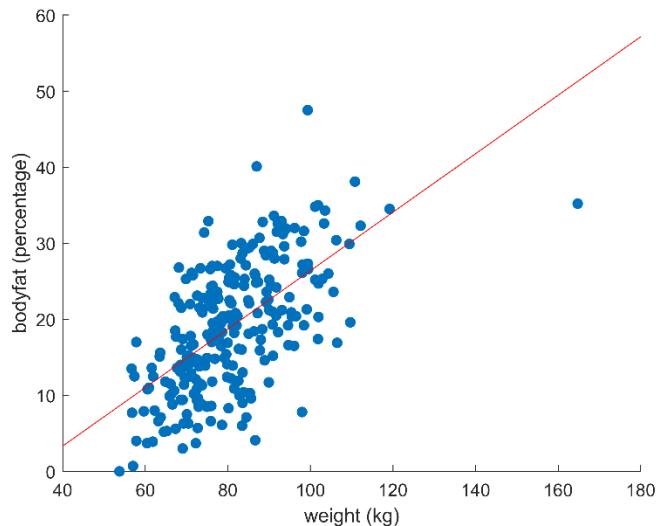
14/02/2022 21:26, by me

0 comments



Interpretation: Caveat

- There is a positive correlation between *weight* and *bodyfat* ($r = 0.61, p < .001$).



Interpretation: Caveat

- There is a positive correlation between *weight* and *bodyfat* ($r = 0.61, p < .001$).
 - *weight* only has a negative coefficient β in the context of *abdomen*, i.e. for fixed *abdomen* predictor
 - a predictor can only be interpreted **in the context** of the other predictors in the model

What means linear?

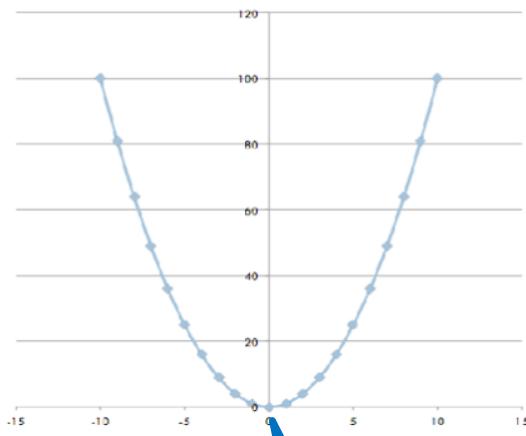
$$y_n = \beta_0 + \beta_1 x_{n,1} + \cdots + \beta_D x_{n,D} + \epsilon_n$$

Linear in the **parameters** -> we can apply arbitrary functions to the raw input variables, e.g.,

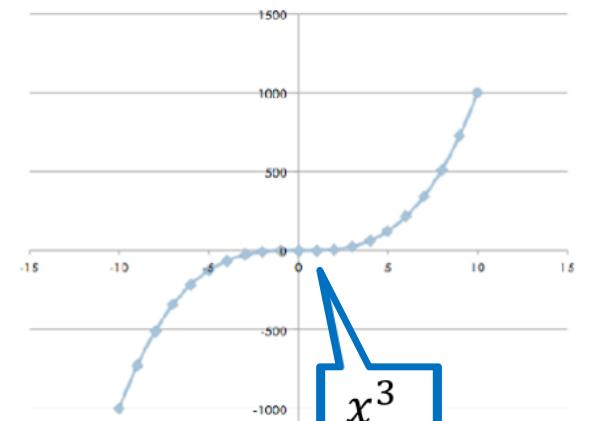
- logarithms, exponentials
- polynomials
- inverse

What means linear

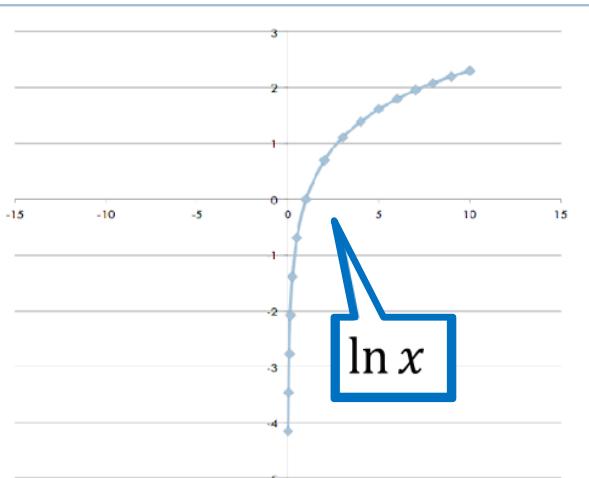
Different transformations



$$x^2$$



$$x^3$$

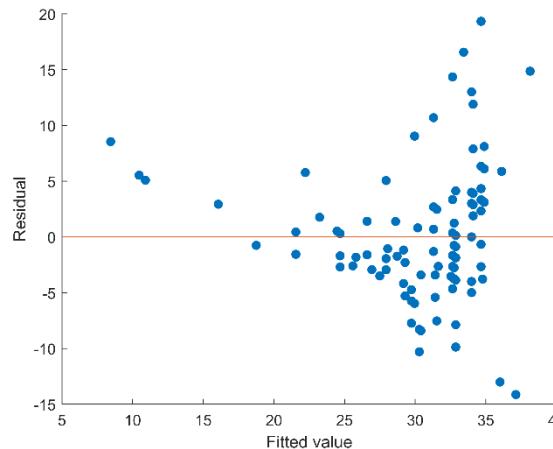
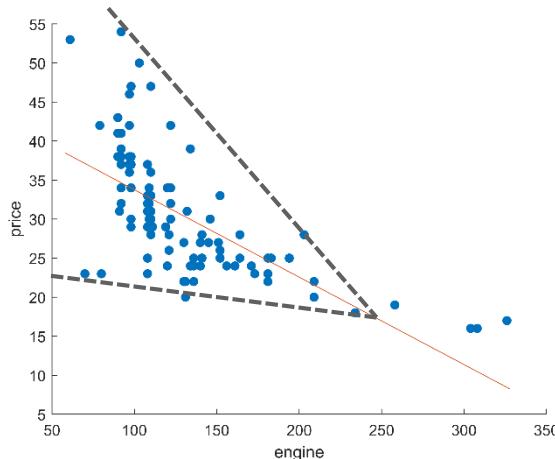


$$\ln x$$

Restrictions of linear models

For some cases, linear regression models are not appropriate:

- the variance of y depends on the mean



Assumption for statistics (t-test, chi, etc.):
 $\epsilon \sim N(0, \sigma^2)$

Restrictions of linear models

For some cases, linear regression models are not appropriate:

- the variance of y depends on the mean
- the range of y is restricted

$$\#bicycles = -2291 + 83 \cdot maxTemp - 13 \cdot minTemp - 890 \cdot precipitation$$

→ prediction \hat{y} can be negative...

Agenda

- Linear Regression
- **Generalized Linear Models**
- Mixed-Effect Models
- Regression for Time-Series Data
- Performance Metrics

Generalized Linear Models

A generalized linear model is composed of a **linear predictor**

$$\pi_n = \beta_0 + \beta_1 x_{n,1} + \cdots + \beta_D x_{n,D}$$

and a **link function**

$$g(\mu_n) = \pi_n$$

with $\mu_n = E[Y|X = x_n]$

Generalized Linear Models

Conditional expectation: the mean μ_n depends on the values of independent variables x_n

is composed of a **linear predictor**

$x_{n,1}$

Each y_n represents the realization of the random variable Y , which is distributed according to a specific probability distribution

and a

function

$$g(\mu_n) = \pi_n$$

$$\text{with } \mu_n = E[Y|X = x_n]$$

Generalized Linear Models

A generalized linear model is composed of a **linear predictor**

$$\pi_n = \beta_0 + \beta_1 x_{n,1} + \dots + \beta_p x_{n,p}$$

In practice (for parameter fitting): observed values y_n are assumed to represent μ_n

and a **link function**

$$g(\mu_n) = \pi_n$$

with $\mu_n = E[Y|X = x_n]$

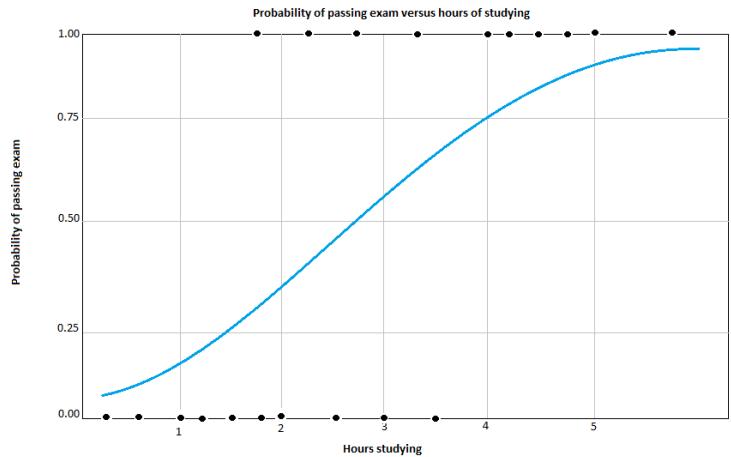
Logistic Regression

In logistic regression, the link function is

$$g(\mu_n) = \log\left(\frac{\mu_n}{1 - \mu_n}\right)$$

and therefore (for fitting)

$$\log\left(\frac{y_n}{1 - y_n}\right) = \beta_0 + \beta_1 x_{n,1} + \cdots + \beta_D x_{n,D}$$



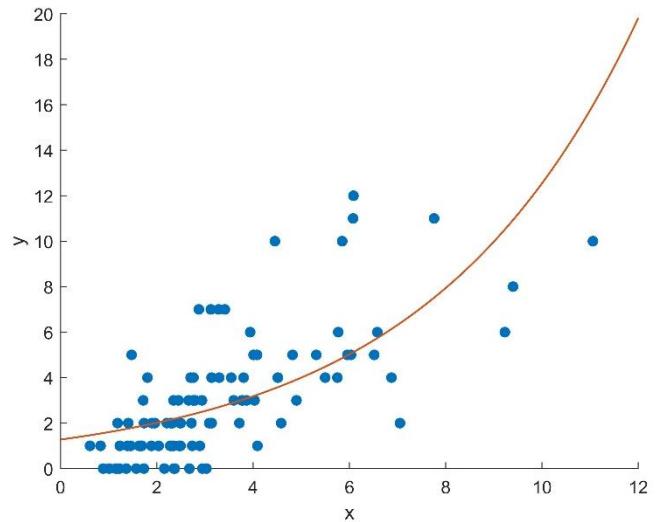
Poisson Regression

In Poisson regression, the link function is

$$g(\mu_n) = \log(\mu_n)$$

and therefore (for fitting)

$$\log(\mu_n) = \beta_0 + \beta_1 x_{n,1} + \cdots + \beta_D x_{n,D}$$



Linear Regression as a special case

For the linear regression, the link function is

$$g(\mu_n) = \mu_n$$

and therefore (for fitting)

$$y_n = \beta_0 + \beta_1 x_{n,1} + \cdots + \beta_D x_{n,D}$$

Example

What type of model would you use for the following tasks?

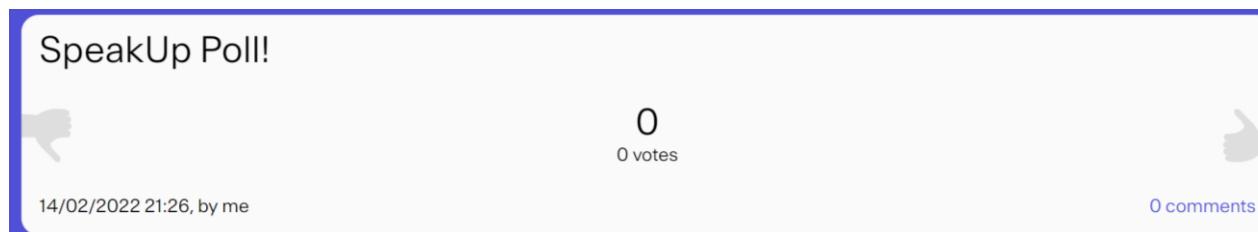
1. Predict the **number of awards** earned by students at one high school.

Predictors include the type of program in which the student was enrolled (e.g., vocational, general or academic) and the score on their final exam in math.

(a) Linear Regression

(b) Logistic Regression

(c) Poisson Regression



Example

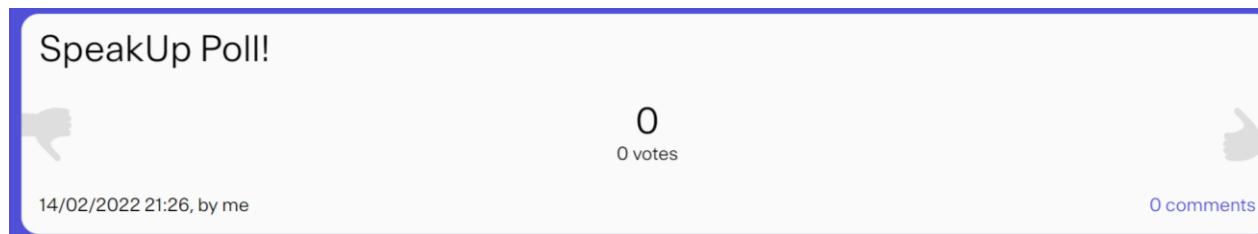
What type of model would you use for the following tasks?

2. Predict whether a student will **solve a task correctly**. Predictors include the difficulty of the task and the number of tasks the student has already solved.

(a) Linear Regression

(b) Logistic Regression

(c) Poisson Regression



Example

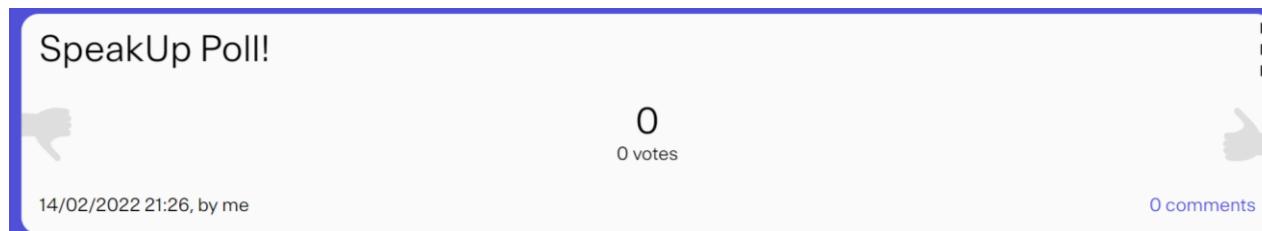
What type of model would you use for the following tasks?

3. Predict the **profit (in \$)** of a company based on their advertising budget on Youtube.

(a) Linear Regression

(b) Logistic Regression

(c) Poisson Regression



Agenda

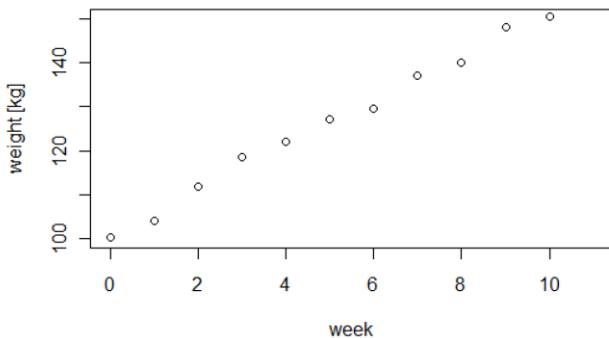
- Linear Regression
- Generalized Linear Models
- **Mixed-Effect Models**
- Regression for Time-Series Data
- Performance Metrics

Why mixed-effect models?

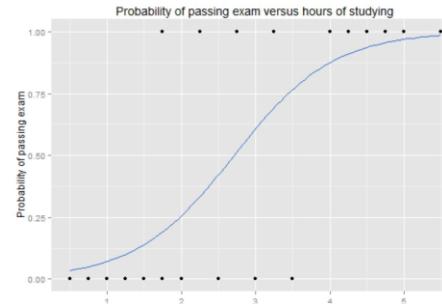
- Useful when we are dealing with **correlated** samples
 - **Grouping** of subjects (e.g., students within a classroom)
 - **Repeated measurements** on each subject over time (e.g., student in flipped classroom course over 10 weeks)

Generalized Linear Models

- Example 1: strength gain by weight training
- Example 2: probability of passing exam of a course c depending on the hours studied



$$y_n = \beta_0 + \beta_1 x_{n,1}$$

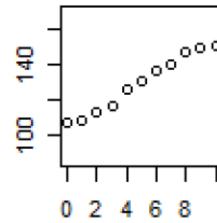
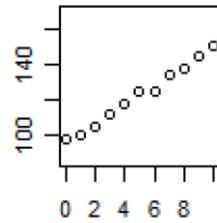
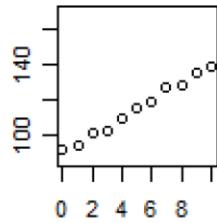


$$\log\left(\frac{y_n}{1 - y_n}\right) = \beta_0 + \beta_1 x_{n,1}$$

“Fixed” Effects

Generalized Linear Mixed Effects Model

- Example 1: strength gain by weight training
 - Each person has individual starting strength



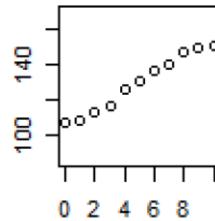
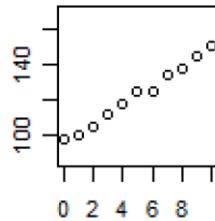
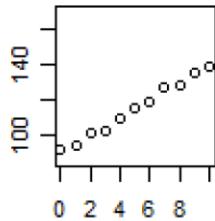
$$y_n = \beta_0 + u_n + \beta_1 x_{n,1} \quad u_n \sim N(0, \sigma_u^2)$$

“Fixed” Effects

“Random” Effect

Generalized Linear Mixed Effects Model

- Example 1: strength gain by weight training
 - Each person has individual starting strength



$$y_n = \beta_0 + u_n + \beta_1 x_{n,1} \quad u_n \sim N(0, \sigma_u^2)$$

Fitting the parameters:

- Fixed effects only: linear least squares
- Mixed effects: maximum likelihood estimation

“Fixed” Effects

“Random” Effect

“Mixed” Effects

Generalized Linear Mixed Effects Model

- In our case, students come from different origins and we assume that students from the same origin are more similar (same education system)
- We therefore use origin (*category*) as a proxy for prior knowledge and add a random intercept to the model

$$passed \sim 1 | category + percentage_correct$$

Notation and code

Intercept	Slope	Notation	Code
Fixed	Fixed	$y_{n,j} = \beta_0 + \beta_1 x_{n,j} + \varepsilon_{n,j}$	y~x
Random	Fixed	$y_{n,j} = (\beta_0 + \mu_{0,j}) + \beta_1 x_{n,j} + \varepsilon_{n,j}$	y~1 group + x
Fixed	Random	$y_{n,j} = \beta_0 + (\beta_1 + \mu_{1,j}) x_{n,j} + \varepsilon_{n,j}$	y~x + (0 + x group)
Random	Random	$y_{n,j} = (\beta_0 + \mu_{0,j}) + (\beta_1 + \mu_{1,j}) x_{n,j} + \varepsilon_{n,j}$	y~(1 group) + x + (0 + x group)

where $y_{n,j}$ denotes the observation of student n from group j

$$\mu_{0,j} \sim \mathcal{N}(0, \sigma_0^2)$$

$$\mu_{1,j} \sim \mathcal{N}(0, \sigma_1^2)$$

$$\text{cov}(\mu_{0,j}, \mu_{1,j}) = 0$$

Movie Madness – Part 1

You are a data scientist working on a movie streaming service. The service wants to create user profiles based on the movies that the users have watched so far. Your colleague has decided to train a regression model to predict the number of movies a user will watch per week. He suggests using one of the following two models, where x_1 denotes the *total time watching movies per day* and x_2 the *chips eaten per day in grams*:

```
mod_1 <- glm(num_movies~1+x1+x2, data=df, family=gaussian)
```

```
mod_2 <- glm(num_movies~1+x1+x2, data=df, family=poisson)
```

- a) How will the outputs of the two suggested models differ?
- b) Which model would you use and why?

Movie Madness – Part 2

When investigating the data further, you realize that male users seem to generally watch more movies than female users. You also observe that the amount of rain (in mm) influences movie consumption, having a much larger effect on females than on males. Adjust the regression model to incorporate the new information (provide an equation using pymer4 notation) and justify your adjustments.

Hint: use `mod_new <- glmer(num_movies~...)`

Summary

- Linear regression is a useful framework for interpreting data and making predictions
- Caveat: be careful when interpreting the models
- Linear regression is flexible, i.e. arbitrary functions can be applied to the raw input data
- Generalized linear models are a more general framework appropriate for response variables from exponential family distributions
- Mixed models allow for capturing correlation in the data
- Modeling time series data requires some type of aggregation