

Recurrent Neural Networks

Machine Learning for Behavioral Data

March 24, 2026

Today's Topic

Week	Lecture/Lab
1	Data Exploration
2	Regression
3	Classification + Model Evaluation
4	Evaluation + Knowledge Tracing
5	Knowledge Tracing
6	Time Series Prediction
7	Unsupervised Learning
8	Spring Break

Getting ready for today's lecture...

➤ Join us on SpeakUp


<https://go.epfl.ch/mlbd-speakup>



Short quiz about the past...

This knowledge tracing model uses the numbers of opportunities the student had per skill as input feature for predicting success.


- a) Additive Factors Model (AFM)
- b) Performance Factors Analysis (PFA)
- c) Bayesian Knowledge Tracing (BKT)




<https://go.epfl.ch/mlbd-speakup>

Short quiz about the past...

Which BKT parameter has been changed between the left and right plot (exactly one)?
The corresponding observations are: [1,1,0,0,1,1,1,0,0,1,1]



$p_0 = 0.6,$
 $p_s = 0.1,$
 $p_g = 0.2,$
 $p_l = 0.3,$
 $p_f = 0.3$



- a) p_g (guess probability)
- c) p_0 (initial probability)

- b) p_l (probability of learning)
- d) p_f (forget probability)

<https://go.epfl.ch/mlbd-speakup>

SpeakUp Poll!

0
0 votes


14/02/2022 21:26, by me

0 comments

Knowledge Tracing – Predicting Future Performance



Subtraction 0-100




Today – Recurrent Neural Networks

- **Deep Knowledge Tracing**
- Parameters and hyperparameter tuning
- Different architectures
- Different tasks:
 - “Many-to-many” versus “Many-to-one”
 - Classification versus Regression

Neural Networks


- Neural networks are able to represent non-linear functions, i.e. $y_n \approx f(\mathbf{x}_n)$ can be non-linear
- Neural networks are able to *learn* the features and the weights (parameters) from the data
- Tutorial: <https://go.epfl.ch/tutorial-nn>

Recurrent Neural Network




$$\mathbf{h}_t^{(1)} = \phi_h (W_{hx} \mathbf{x}_t + W_{h^{(1)}h^{(1)}}^{(1)} \mathbf{h}_{t-1}^{(1)} + \mathbf{b}^{(1)})$$


Recurrent Neural Network




Recurrent Neural Network




Deep Knowledge Tracing




Deep Knowledge Tracing




Deep Knowledge Tracing




Deep Knowledge Tracing




Deep Knowledge Tracing – Computational Graph




Deep Knowledge Tracing – Computational Graph




Deep Knowledge Tracing – Computational Graph




Deep Knowledge Tracing – Computational Graph




Deep Knowledge Tracing – Computational Graph



Deep Knowledge Tracing – Computational Graph



Training a DKT model: Binary Crossentropy Loss



Training and Prediction using DKT

- Training: gradient descent
- Prediction: compute inference in the network (see computational graph)

Today – Recurrent Neural Networks

- Deep Knowledge Tracing
- **Parameters and hyperparameter tuning**
- Different architectures
- Different tasks:
 - “Many-to-many” versus “Many-to-one”
 - Classification versus Regression

RNNs – Specifying Parameters

```
[ ] # Specify the model hyperparameters. Full descriptions included in the demo notebook!  
params = {}  
  
params['batch_size'] = 32  
params['mask_value'] = -1.0  
params['verbose'] = 1  
params['best_model_weights'] = 'weights/bestmodel'  
params['optimizer'] = 'adam'  
params['recurrent_units'] = 16  
params['epochs'] = 20  
params['dropout_rate'] = 0.1
```

RNNs – Tuning hyperparameters

```
[ ] # Specify the model hyperparameters. Full descriptions included in the demo notebook!  
params = {}  
  
params['batch_size'] = 32  
params['mask_value'] = -1.0  
params['verbose'] = 1  
params['best_model_weights'] = 'weights/bestmodel'  
params['optimizer'] = 'adam'  
params['recurrent_units'] = 16  
params['epochs'] = 20  
params['dropout_rate'] = 0.1
```


RNNs – Tuning hyperparameters

- Optimal number of epochs can be found using callbacks
- Other parameters can be tuned using for example:
 - a) Train-Validation-Test split
 - b) Train-Test split, using a k-fold cross validation on the training data to determine the optimal parameters


Today – Recurrent Neural Networks

- Parameters and hyperparameter tuning
- **Different architectures**
- Different tasks:
 - “Many-to-many” versus “Many-to-one”
 - Classification versus Regression

Recurrent Neural Network




Vanilla RNN - revisited




$$\mathbf{h}_t = \tanh(W_{hx}\mathbf{x}_t + W_{hh}\mathbf{h}_{t-1} + \mathbf{b}_h)$$

Long-Short Term Memory Network (LSTM)



- Two states:
 - Hidden state h_{t-1}
 - Cell state c_{t-1}

Long-Short Term Memory Network (LSTM)




- ① Updating the gates:
- f forget gate: whether to erase cell
 - i input gate: whether to write to cell
 - o output gate: how much to reveal cell

$$f_t = \sigma(W_{fx}x_t + W_{fh}h_{t-1} + b_f)$$

$$i_t = \sigma(W_{ix}x_t + W_{ih}h_{t-1} + b_i)$$

$$o_t = \sigma(W_{ox}x_t + W_{oh}h_{t-1} + b_o)$$

Long-Short Term Memory Network (LSTM)




② Updating the states:

$$c_t = f_t \odot c_{t-1} + i_t \odot \tanh(W_{cx}x_t + W_{ch}h_{t-1} + b_c)$$


$$h_t = o_t \odot \tanh(c_t)$$

Gated Recurrent Units (GRU)



- Only one state (got rid of cell):
 - Hidden state h_{t-1}

Gated Recurrent Units (GRU)




① Updating the gates:

- r reset gate: how much of the previous state to remember
- u update gate: how much of the new state is just a copy of the old state

$$r_t = \sigma(W_{rx}x_t + W_{th}h_{t-1} + b_r)$$

$$u_t = \sigma(W_{ux}x_t + W_{uh}h_{t-1} + b_u)$$


Gated Recurrent Units (GRU)



② Get candidate hidden state:

$$\tilde{h}_t = \tanh(W_{hx}x_t + W_{ht}(r_t \odot h_{t-1}) + b_h)$$


Gated Recurrent Units (GRU)



③ Updating the state:

$$h_t = u_t \odot h_{t-1} + (1 - u_t) \odot \tilde{h}_t$$

Same input/output – different architectures




Today – Recurrent Neural Networks


- Deep Knowledge Tracing
- Parameters and hyperparameter tuning
- Different architectures
- **Different tasks:**
 - “Many-to-many” versus “Many-to-one”
 - Classification versus Regression

Many-to-many aka the Tracing Task

- Prediction of a target variable o_t at each time step t




Computational Graph – Many-to-many




Many-to-one aka the Time-Series Prediction Task

- Prediction of a target variable o after $t \leq T$ time steps, where T is the total number of time steps




Computational Graph – Many-to-one




Today – Recurrent Neural Networks

- Deep Knowledge Tracing
- Parameters and hyperparameter tuning
- Different architectures
- **Different tasks:**
 - “Many-to-many” versus “Many-to-one”
 - **Classification versus Regression**


Classification vs. Regression



Classification vs. Regression: Output Layer



Classification vs. Regression: Training Loss



Your Turn

- Given:
 - Data from a MOOC
 - An LSTM for predicting quiz performance of a student for every week of the course (tracing task)
- Your Task:
 - 1) Adjust the `create_model` function in order to predict pass/fail after 5 weeks of the course (time series prediction task) and send us the binary accuracy + AUC
 - Hint 1: `return_sequences=False`
 - Hint 2: what does `TimeDistributed(...)` do?
 - 2) Tune hyperparameters of your choice and send us binary accuracy and AUC

Summary

- Deep Knowledge Tracing
- Parameters and hyperparameter tuning
- Different architectures
- Different tasks:
 - “Many-to-many” versus “Many-to-one”
 - Classification versus Regression